

Code by Sarah Wiegreffe (saw@gatech.edu)
Fall 2019

```
import numpy as np
```

```
import torch
from torch import nn
import random
```

Do not modify these imports.

```
class ClassificationTransformer(nn.Module):
```

"""

A single-layer Transformer which encodes a sequence of text and performs binary classification.

The model has a vocab size of V , works on sequences of length T , has an hidden dimension of H , uses word vectors also of dimension H , and operates on minibatches of size N .

"""

```
def __init__(self, word_to_ix, hidden_dim=128, num_heads=2, dim_feedforward=2048, dim_k=96, dim_v=96, dim_q=96, max_length=43):
```

"""
:param word_to_ix: dictionary mapping words to unique indices
:param hidden_dim: the dimensionality of the output embeddings that go into the final layer
:param num_heads: the number of Transformer heads to use
:param dim_feedforward: the dimension of the feedforward network model
:param dim_k: the dimensionality of the key vectors
:param dim_q: the dimensionality of the query vectors
:param dim_v: the dimensionality of the value vectors
"""

```
super(ClassificationTransformer, self).__init__()
```

```
assert hidden_dim % num_heads == 0
```

```
self.num_heads = num_heads
self.word_embedding_dim = hidden_dim
self.hidden_dim = hidden_dim
self.dim_feedforward = dim_feedforward
self.max_length = max_length
self.vocab_size = len(word_to_ix)
```

```
self.dim_k = dim_k
self.dim_v = dim_v
self.dim_q = dim_q
```

```
seed_torch(0)
```

```
#####
# Deliverable 1: Initialize what you need for the embedding lookup.      #
# You will need to use the max_length parameter above.                  #
# This should take 1-2 lines.                                           #
# Initialize the word embeddings before the positional encodings.       #
# Don't worry about sine/cosine encodings- use positional encodings.    #
#####
self.token_embedding = torch.nn.Embedding(self.vocab_size, self.word_embedding_dim)
self.positional_encodings = torch.nn.Embedding(self.max_length, self.word_embedding_dim)

#####
#                               END OF YOUR CODE                          #
#####
```

```
#####
# Deliverable 2: Initializations for multi-head self-attention.         #
# You don't need to do anything here. Do not modify this code.        #
#####
```

```

# Head #1
self.k1 = nn.Linear(self.hidden_dim, self.dim_k)
self.v1 = nn.Linear(self.hidden_dim, self.dim_v)
self.q1 = nn.Linear(self.hidden_dim, self.dim_q)

# Head #2
self.k2 = nn.Linear(self.hidden_dim, self.dim_k)
self.v2 = nn.Linear(self.hidden_dim, self.dim_v)
self.q2 = nn.Linear(self.hidden_dim, self.dim_q)

self.softmax = nn.Softmax(dim=-2)
self.attention_head_projection = nn.Linear(self.dim_v * self.num_heads, self.hidden_dim)
self.norm_mh = nn.LayerNorm(self.hidden_dim)

#####
# Deliverable 3: Initialize what you need for the feed-forward layer.      #
# Don't forget the layer normalization.                                   #
#####
self.linear1 = nn.Linear(self.hidden_dim, self.dim_feedforward)
self.linear2 = nn.Linear(self.dim_feedforward, self.hidden_dim)
self.norm_mh_2 = nn.LayerNorm(self.hidden_dim)
#####
#                               END OF YOUR CODE                               #
#####

#####
# Deliverable 4: Initialize what you need for the final layer (1-2 lines). #
#####
self.final_linear = nn.Linear(self.hidden_dim, 1)
self.sigmoid = nn.Sigmoid()
#####
#                               END OF YOUR CODE                               #
#####

```

```

def forward(self, inputs):
    """
    This function computes the full Transformer forward pass.
    Put together all of the layers you've developed in the correct order.

    :param inputs: a PyTorch tensor of shape (N,T). These are integer lookups.

    :returns: the model outputs. Should be normalized scores of shape (N,1).
    """
    outputs = None
    #####
    # Deliverable 5: Implement the full Transformer stack for the forward pass. #
    # You will need to use all of the methods you have previously defined above. #
    # You should only be calling ClassificationTransformer class methods here. #
    #####
    outputs = self.final_layer(self.feedforward_layer(self.multi_head_attention(self.embed(inputs))))
    #####
    #                               END OF YOUR CODE                               #
    #####
    return outputs

```

```

def embed(self, inputs):
    """
    :param inputs: intTensor of shape (N,T)
    :returns embeddings: floatTensor of shape (N,T,H)
    """
    embeddings = None
    #####
    # Deliverable 1: Implement the embedding lookup. #

```

```

# Note: word_to_ix has keys from 0 to self.vocab_size - 1
# This will take a few lines.
#####
tokens = self.token_embedding(inputs)
n, t, w = tokens.size()
positions = self.positional_encodings(torch.arange(t))[None, :, :].expand(n, t, w)
embeddings = tokens + positions
#####
#                               END OF YOUR CODE                               #
#####
return embeddings

```

```

def multi_head_attention(self, inputs):

```

```

    """

```

```

    :param inputs: float32 Tensor of shape (N,T,H)

```

```

    :returns outputs: float32 Tensor of shape (N,T,H)

```

```

    Traditionally we'd include a padding mask here, so that pads are ignored.

```

```

    This is a simplified implementation.

```

```

    """

```

```

    outputs = None

```

```

    #####

```

```

    # Deliverable 2: Implement multi-head self-attention followed by add + norm. #

```

```

    # Use the provided 'Deliverable 2' layers initialized in the constructor. #

```

```

    #####

```

```

    h1 =

```

```

    torch.matmul(self.softmax(torch.matmul(self.q1(inputs), self.k1(inputs).transpose(-2, -1))/np.sqrt(self.dim_k)), self.v1(inputs))

```

```

    h2 =

```

```

    torch.matmul(self.softmax(torch.matmul(self.q2(inputs), self.k2(inputs).transpose(-2, -1))/np.sqrt(self.dim_k)), self.v2(inputs))

```

```

    mh = self.attention_head_projection(torch.cat((h1, h2), axis=2))

```

```

    outputs = self.norm_mh(mh+inputs)

```

```

    #####

```

```

    #                               END OF YOUR CODE                               #

```

```

    #####

```

```

    return outputs

```

```

def feedforward_layer(self, inputs):

```

```

    """

```

```

    :param inputs: float32 Tensor of shape (N,T,H)

```

```

    :returns outputs: float32 Tensor of shape (N,T,H)

```

```

    """

```

```

    outputs = None

```

```

    #####

```

```

    # Deliverable 3: Implement the feedforward layer followed by add + norm. #

```

```

    # Use a ReLU activation and apply the linear layers in the order you #

```

```

    # initialized them. #

```

```

    # This should not take more than 3-5 lines of code. #

```

```

    #####

```

```

    outputs = self.norm_mh_2(inputs + self.linear2(nn.functional.relu(self.linear1(inputs))))

```

```

    #####

```

```

    #                               END OF YOUR CODE                               #

```

```

    #####

```

```

    return outputs

```

```

def final_layer(self, inputs):

```

```

    """

```

```

    :param inputs: float32 Tensor of shape (N,T,H)

```

```

    :returns outputs: float32 Tensor of shape (N,1)

```

```

    """

```

```

    outputs = None

```

```

    #####

```

```

    # Deliverable 4: Implement the final layer for the Transformer classifier. #

```

```

    # This should not take more than 2 lines of code. #

```

```
#####
outputs = self.sigmoid(self.final_linear(inputs[:,0,:]))
#####
#                               #
#           END OF YOUR CODE           #
#                               #
#####
return outputs
```

```
def seed_torch(seed=0):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed(seed)
    torch.backends.cudnn.benchmark = False
    torch.backends.cudnn.deterministic = True
```